

SIDDESH KAMBLE

Syracuse, NY | (315) 450-9024 | siddesh.kamble@techsmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

SUMMARY

Software Engineer with 5+ years of experience creating production-scale Python services, LLM pipelines, and retrieval infrastructure across enterprise, product & fintech platforms. High-concurrency backend systems, conversational AI platforms serving 50K+ users, and low-latency data services powering intelligent automation and data-driven workflows across distributed cloud environments. Multi-agent reasoning architectures using LangGraph, RAG pipelines, LangChain, FAISS, and OpenAI APIs applied to large financial document corpora for compliance validation and enterprise knowledge retrieval.

EXPERIENCE

STRIPE Aug 2025 – Current
Software Engineer, Machine Learning (Backend Systems) – AI-Assisted Compliance & Document Intelligence Platform

- Designed **LLM-powered document processing pipelines using Python, LangChain, and OpenAI APIs**, extracting structured insights from **120K+ financial and regulatory documents** and reducing manual analyst review time by **40%**.
- Architected **retrieval infrastructure for enterprise knowledge search** using **vector embeddings, FAISS indexing, and metadata filtering**, enabling low-latency contextual retrieval across **15M+ internal records** and improving response grounding accuracy by **30%**.
- Developed **backend inference services using Python, FastAPI, and PostgreSQL**, introducing **Redis caching and async processing** to sustain high-concurrency workloads while maintaining **sub-250ms response latency**.
- Engineered **agent-orchestrated reasoning workflows using LangGraph**, enabling automated document validation, multi-step compliance checks, and structured tool execution across enterprise data pipelines.
- Deployed **containerized AI services with Docker and cloud infrastructure (AWS/GCP)** and implemented **CI/CD deployment pipelines**, reducing release cycles from weekly to **same-day production deployments**.
- Built **LLM evaluation and reliability tooling** using automated prompt testing datasets and telemetry monitoring, improving response stability and reducing hallucination risk in production workflows.

QUANTIPHI Jun 2021 - Oct 2023
Software Engineer – Conversational AI

- Implemented production conversational platforms using **Dialogflow ES/CX, Python, and REST APIs**, supporting **50K+ users** with **99–99.5% uptime** while automating high-volume customer support workflows.
- Established backend conversation orchestration services using **Python, Flask, and microservice APIs**, improving intent routing accuracy and reducing human-agent escalation across enterprise support channels.
- Integrated conversational interfaces with enterprise data services through **Python, REST APIs, and AWS infrastructure**, enabling secure real-time data retrieval and automated workflow execution.
- Optimized dialogue performance through conversation analytics pipelines using **Python, SQL, and A/B testing**, identifying drop-off patterns and improving task completion rates across high-traffic intents.
- Facilitated internal engineering quality through **code reviews, CI/CD pipelines**, and shared development tooling using **Docker, Git, and Jenkins**, improving deployment stability and onboarding efficiency across engineering teams.
- Coordinated with product managers and backend engineers to deliver scalable conversational workflows using **Dialogflow CX, Python APIs, and AWS services**.

SPACE INFOLAB Jan 2019 - May 2021
Software Developer – Backend

- Designed backend application services using **Python** and **Django**, implementing modular service layers for authentication, request validation, and data access while supporting high-volume internal workflow operations.
- Implemented asynchronous job processing frameworks using **Python, Celery, and Redis**, enabling reliable background execution for compute-intensive tasks and improving throughput for large data-processing workloads.
- Optimized relational data models and complex queries using **PostgreSQL, SQL indexing**, and query profiling, significantly improving response times across analytics and reporting services.
- Developed automated data ingestion and transformation pipelines using **Python, Pandas**, and batch-processing workflows to validate, clean, and structure operational datasets for downstream services.
- Managed containerized backend environments using **Docker, Linux**, and application monitoring tools, stabilizing service deployments and streamlining development-to-production workflows.

TECHNICAL SKILLS

Programming: Python, SQL, JavaScript
Backend Development: FastAPI, Django, Flask, REST API Design, Microservices Architecture, Asynchronous Processing
AI / ML Systems: LLM Application Development, Retrieval-Augmented Generation (RAG), LangChain, LangGraph, Vector Embeddings, Prompt Engineering, AI Workflow Automation, Machine Learning, AgenticAI
Conversational AI: Dialogflow ES, Dialogflow CX, Intent Routing, Dialogue Optimization, Conversational Workflow Design
Data Engineering: Data Pipelines, ETL, Pandas, SQL Query Optimization, Relational Data Modeling, DataValidation
Databases: PostgreSQL, Redis, SQL Indexing, Query Profiling, Oracle, MongoDB, DynamoDB
Cloud & Infrastructure: AWS, Google Cloud Platform (GCP), Containerized Services, Distributed Systems
DevOps & Tools: Docker, CI/CD Pipelines, Git, Jenkins, Linux, Application Monitoring
AI Reliability: LLM Evaluation Pipelines, Prompt Testing, Model Output Validation, Telemetry Monitoring

CERTIFICATIONS

[Google CCAI](#)
[AWS Certified Developer – Associate](#)
[Oracle Cloud Infrastructure 2025 Certified Generative AI Professional](#)

EDUCATION

Syracuse University, Syracuse, NY | Master of Science in Computer Science | Dec 2025